

Kubernetes

Going 0 to 100

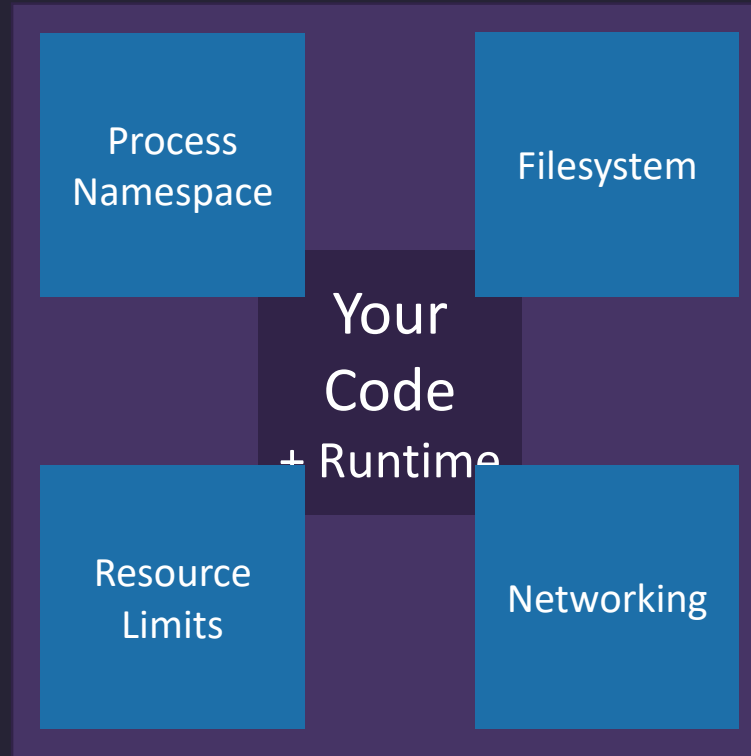
Scott Holden



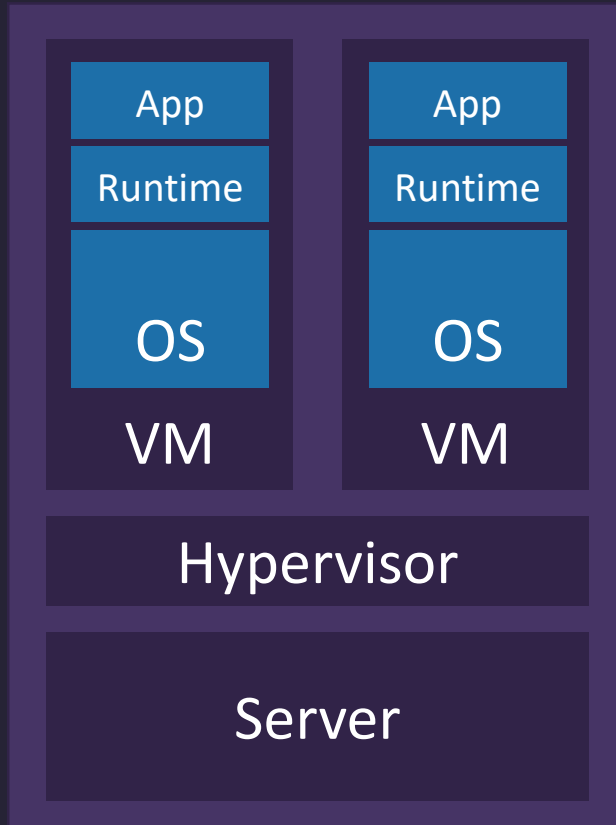
Let's
Revisit
Containers



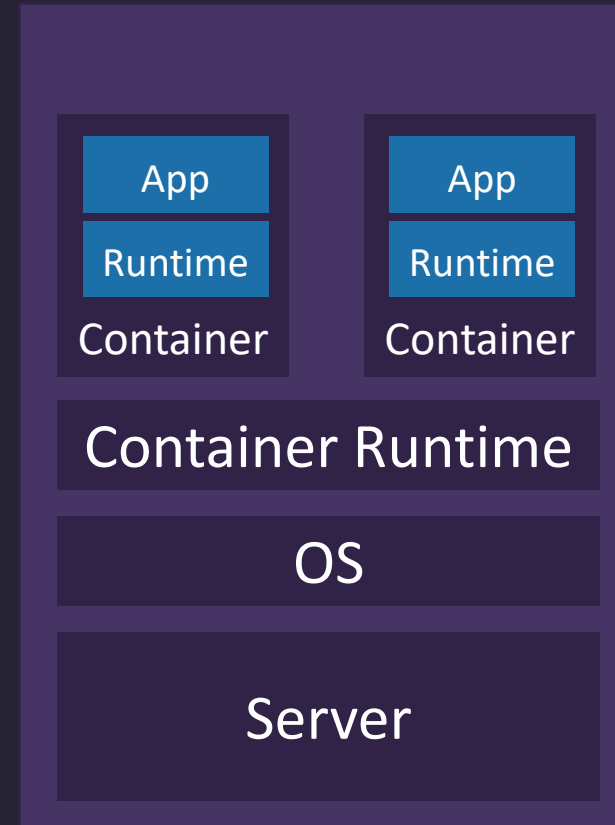
Sandbox →



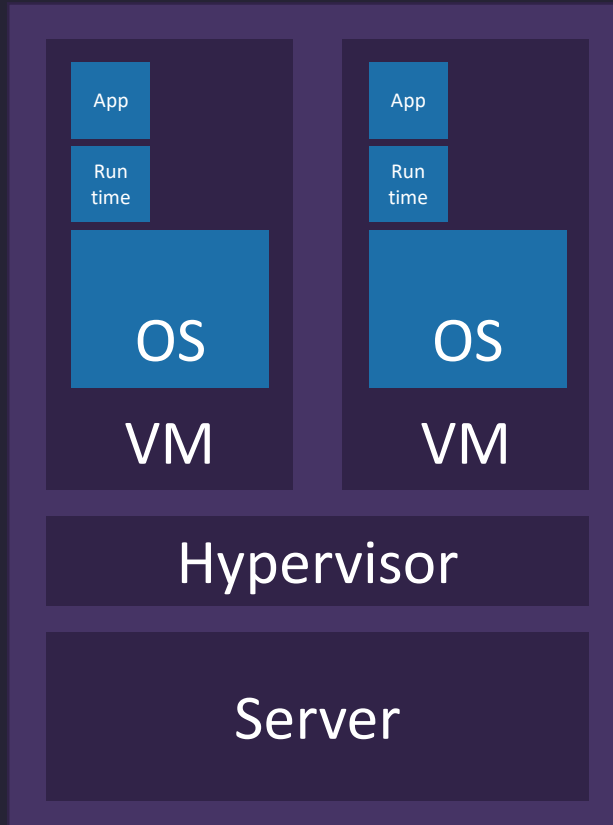
Container



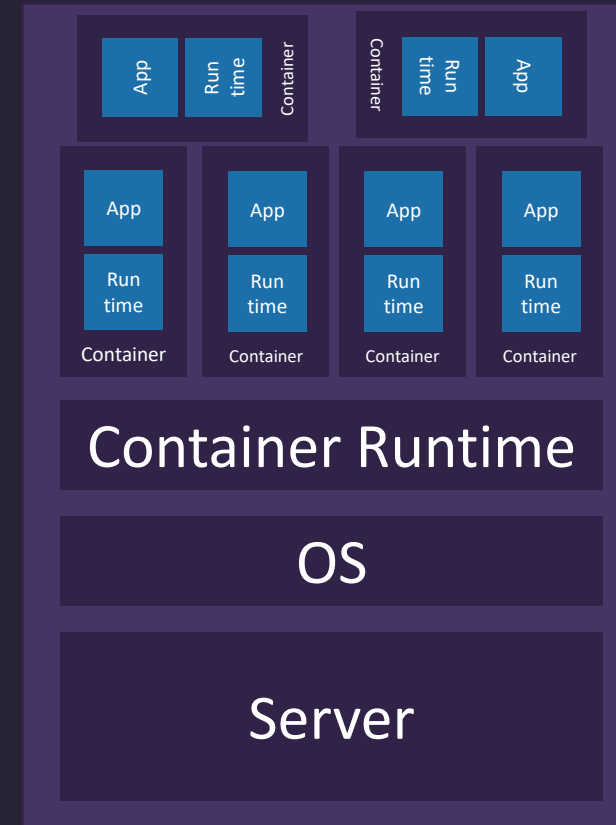
Virtual Machine



Container



Virtual Machine



Container

**Container
Image**
+ Configuration



**Container
Runtime**



**Running
Container**



docker

Demo:

Running A Simple Container


Container Image

Container
Image

Container Image

Container Image

scratch







Custom Container
Image

scratch



Demo:

Making A Container Image

Docker
Hub

Private
Container
Registry

Pull

Push

Pull

Custom Container
Image

Nginx

Alpine

scratch



What's The
Challenge?

One container, one host

Manageable

A few containers, a few hosts

Getting tricky

Lots of containers, many hosts

More than a full time job!

One host fails

Move the container

Containers move

Service discovery

Is this container healthy?

Monitoring & recovery

Scaling in & out

Which host?

Deploying new versions

How to update?

Required configuration

Check every host

Introducing Kubernetes



Scheduling



Affinity/anti-affinity



Health monitoring



Failover



Scaling



Networking



Service discovery



Coordinated
app
upgrades

Kubernetes is a Scheduler*

*In terms of running containers

Kubernetes

Masters & Nodes

Master (Node)

Runs Kubernetes

Made up of a few key services

Talks with the worker nodes

Normally a cluster

100+

kube-apiserver

(Front end, what kubectl talks to)

kube-scheduler

(Decides what node gets what)

etcd*

(Internal backing store)

kube-controller-manager

(Runs the internal controllers)

Master Node

cloud-controller-manager

(Underlying Cloud provider comms)

(Worker) Node

Some form of compute

Runs your workloads

Container runtime lives here

100+

kubelet

(Controls what containers to run)

kube-proxy


(Sets up and maintains networking)

Container Runtime

(Docker, etc, This is what runs the containers)


Worker Node

Master Nodes



Control Plane

Worker Nodes



IaaS

PaaS

Kubernetes

Master Node
Management

Application Level
Deployment

Worker Node
Management

Workloads
Abstracted From
Hardware

Cloud Managed Masters

Master Nodes 'as a Service'

Azure Kubernetes Service

Google Container Engine

IaaS

PaaS

Kubernetes

Master Node
Management

Application Level
Deployment

Worker Node
Management

Workloads
Abstracted From
Hardware

IaaS

PaaS

Kubernetes

Worker Node
Management

Application Level
Deployment

Workloads
Abstracted From
Hardware

Master Node
Management

Kubernetes

Nuts & Bolts

Pod

Lowest component

Made up of one or more containers

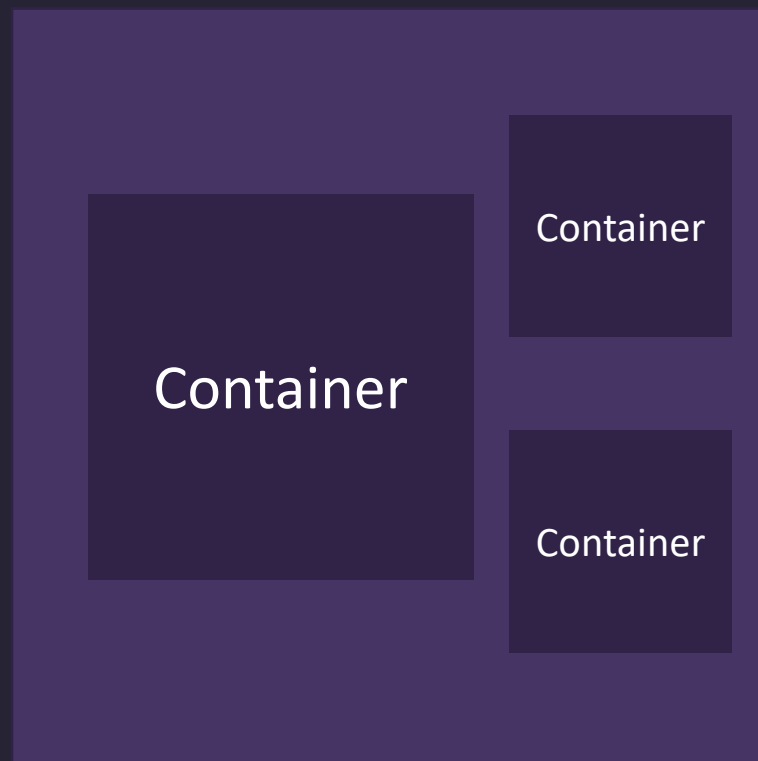
Single instance



A diagram illustrating the relationship between a Pod and a Container. It consists of a large, light purple square representing the Pod, which contains a smaller, dark purple square representing the Container. The word "Container" is written in white text inside the dark purple square.

Container

Pod



Pod



A diagram illustrating the relationship between a Pod and a Container. It consists of a large, light purple square representing the Pod, which contains a smaller, dark purple square representing the Container. The word "Container" is written in white text inside the dark purple square.

Container

Pod

Demo:


Single Pod

ReplicaSet

Sole purpose to maintain a set of pods

Templated definition of a pod

ReplicaSet



replicas: 2

Container

Pod

Container

Pod

ReplicaSet

Container

Pod

replicas: 3

Container

Pod


Container

Pod

Container

Pod

ReplicaSet



replicas: 2


Container

Pod

Container

Pod

ReplicaSet



replicas: 2


Container

Pod

Container

Pod

ReplicaSet



replicas: 2


Container

Pod

Container

Pod

ReplicaSet



replicas: 1

Container

Pod

Demo:

ReplicaSet

100+

Other Types

ReplicaSet – Replicas of a Pod

Job – Run & terminate


StatefulSet – Sticky identity for Pods

DaemonSet – Linked to Node lifecycle

Deployment


Level above a ReplicaSet
Designed to handle updates
via multiple ReplicaSet's
Deploy & Rollback

Deployment



replicas: 2

ReplicaSet



replicas: 2


Container
v1

Pod

Container
v1


Pod

Deployment




replicas: 2

ReplicaSet


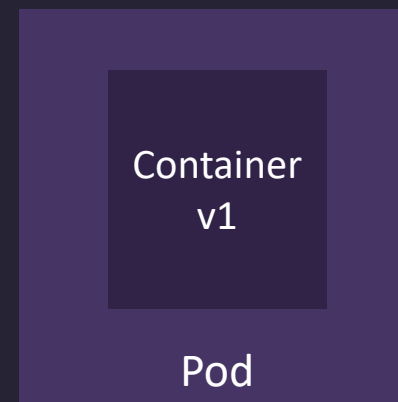
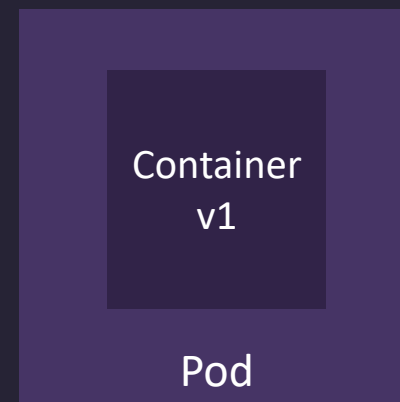


replicas: 2


ReplicaSet



replicas: 1




Deployment




replicas: 2

ReplicaSet







replicas: 0

ReplicaSet



replicas: 2



Demo: Deployment

Service

Expose Pod's as a network service

Simple DNS names

Uses an underlying kube-proxy*

Service

Port 80
(http)

Deployment

Container
v1

Pod

replicas: 2

ReplicaSet

Container
v1

Pod

replicas: 2

Container
v1

Pod

Container
v1

Pod

Service Publishing

ClusterIP (Default) – Internal only

NodePort – Expose a port on every node

LoadBalancer – Use an external LB

Demo:

Service & Publishing

100+

Ingress

Exposes HTTP & HTTPS routes

Reverse Proxy

Single IP, multiple services

Requires an Ingress Controller

mycluster.com

Ingress

/foo/

-> ServiceA

/bar/

-> ServiceB

Service A

Port 80

(http)

Service B

Port 80

(http)

Cont

Container
A

Pod

Pod

Cont

Container
B

Pod

Pod

100+

Demo:

Ingress via Nginx

Kubernetes Extras & More

100+

Autoscaling

Scale Pods via Horizontal Pod Autoscaler
Control Replicas via Metrics

Scaling Nodes is provider specific

100+

Volumes

Directory on a Disk or another Container

Lots of backing providers!

PersistentVolume's give more control
around lifecycle.

100+

Namespaces

Logical separation of a cluster

Namespace per team

Some objects (nodes) aren't
namespace'd

100+

Affinity and anti-affinity

Greater control over node mapping
Via *nodeSelector*

Even inter-pod affinity!

100+

Taints & Tolerations

Great for multiple worker node pools

Allows for 'soft' or 'hard' requirement

Dedicated nodes or specialised hardware

100+

Resource Quotas

Memory, CPU, Storage
Limits & Requests

Request 1gb, Limit 2gb


Azure Kubernetes Service

Master Nodes



Control Plane


Worker Nodes




Master Nodes



Worker Nodes



Control Plane



Built-in auto scaling




Elastically burst using ACI

Global data centers



Geo-replicated container
registry






Control access through
AAD and RBAC

Advanced networking
via Azure CNI



Safeguard keys and
secrets with Key Vault

SOC, HIPAA, PCI & more



Tips & Tricks

Keep it simple!

Start simple

Don't be too verbose

Easy to read = easy to review

Group where appropriate

Deployment & Service?

Single file!

100+ deployments, services, and more?

Multiple files

Single Pod or ReplicaSet?

No one likes a naked pod...

ReplicaSet with 'replicas: 1'

Even better, go for a deployment!

Avoid host bound networking

HostNetwork, HostIP, HostPort
Prefer Services and/or Ingress

Use labels!

Name is a great start

Version is helpful

Part-Of for grouping

Component for identification

Container image tags

Don't use :latest

Be specific with versions

Image pull behaviour is impacted by this

Right size node

Tiny nodes are great for dev...

But not so great for prod!

Pod scaling requires space (deployment!)

You can horizontally scale nodes!

Kubectl

Learn the CLI

kubectl --help is your friend


Apply everything in a folder

kubectl apply -f

Start stateless

External state is simple

Volumes, while incredibly powerful,
require planning



Thank You!
Questions?